

[Click to Take the FREE NLP Crash-Course](#)

Search...



# How to Use Word Embedding Layers for Deep Learning with Keras

by **Jason Brownlee** on October 4, 2017 in **Deep Learning for Natural Language Processing**

Tweet

Share

Share

Last Updated on October 3, 2019

**Word embeddings** provide a dense representation of words and their relative meanings.

They are an improvement over sparse representations used in simpler bag of word model representations.

Word embeddings can be learned from text data and reused among projects. They can also be learned as part of fitting a neural network on text data.

In this tutorial, you will discover how to use word embeddings for deep learning in Python with Keras.

After completing this tutorial, you will know:

- About word embeddings and that Keras supports word embeddings via the Embedding layer.
- How to learn a word embedding while fitting a neural network.
- How to use a pre-trained word embedding in a neural network.

Discover how to develop deep learning models for text classification, translation, photo captioning and more [in my new book](#), with 30 step-by-step tutorials and full source code.

Let's get started.

- **Update Feb/2018:** Fixed a bug due to a change in the underlying APIs.
- **Updated Oct/2019:** Updated for Keras 2.3 and TensorFlow 2.0.



How to Use Word Embedding Layers for Deep Learning with Keras

How to Use Word Embedding Layers for Deep Learning with Keras

Photo by [thisguy](#), sc

**Start Machine Learning**

# Tutorial Overview

This tutorial is divided into 3 parts; they are:

1. Word Embedding
2. Keras Embedding Layer
3. Example of Learning an Embedding
4. Example of Using Pre-Trained GloVe Embedding

## Need help with Deep Learning for Text Data?

Take my free 7-day email course

Click to sign-up and also get a free

Start Your FREE C

### Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**. Find out how in this *free* and *practical* course.

START MY EMAIL COURSE

## 1. Word Embedding

A word embedding is a class of approaches for representing words and documents using a dense vector representation.

It is an improvement over more the traditional bag-of-word model encoding schemes where large sparse vectors were used to represent each word or to score each word within a vector to represent an entire vocabulary. These representations were sparse because the vocabularies were vast and a given word or document would be represented by a large vector comprised mostly of zero values.

Instead, in an embedding, words are represented by dense vectors where a vector represents the projection of the word into a continuous vector space.

The position of a word within the vector space is learned from text and is based on the words that surround the word when it is used.

The position of a word in the learned vector space is referred to as its embedding.

Two popular examples of methods of learning word embeddings from text include:

- Word2Vec.
- GloVe.

In addition to these carefully designed methods, a word

learning model. This can be a slower approach, but ta

Start Machine Learning

## 2. Keras Embedding Layer

Keras offers an [Embedding](#) layer that can be used for neural networks on text data.

It requires that the input data be integer encoded, so that each word is represented by a unique integer. This data preparation step can be performed using the [Tokenizer API](#) also provided with Keras.

The Embedding layer is initialized with random weights and will learn an embedding for all of the words in the training dataset.

It is a flexible layer that can be used in a variety of ways, such as:

- It can be used alone to learn a word embedding
- It can be used as part of a deep learning model with other layers
- It can be used to load a pre-trained word embedding

The Embedding layer is defined as the first hidden layer in a model.

It must specify 3 arguments:

- **input\_dim**: This is the size of the vocabulary in the input data. If the input is integer encoded, then the size of the vocabulary is the number of unique integers. For example, if the input is integer encoded words from 0 to 199, inclusive, then the size of the vocabulary is 200.
- **output\_dim**: This is the size of the vector space in which words will be embedded. For example, it could be 32 or 100 or even larger. Test different values for your problem.
- **input\_length**: This is the length of input sequences, as you would define for any input layer of a Keras model. For example, if all of your input documents are comprised of 1000 words, this would be 1000.

For example, below we define an Embedding layer with a vocabulary of 200 (e.g. integer encoded words from 0 to 199, inclusive), a vector space of 32 dimensions in which words will be embedded, and input documents that have 50 words each.

```
1 e = Embedding(200, 32, input_length=50)
```

The Embedding layer has weights that are learned. If you save your model to file, this will include weights for the Embedding layer.

The output of the *Embedding* layer is a 2D vector with one embedding for each word in the input sequence of words (input document).

If you wish to connect a *Dense* layer directly to an Embedding layer, you must first flatten the 2D output matrix to a 1D vector using the *Flatten* layer.

Now, let's see how we can use an Embedding layer in practice.

## 3. Example of Learning an Emb

Start Machine Learning

In this section, we will look at how we can learn a word embedding while fitting a neural network on a [text classification](#) problem.

We will define a small problem where we have 10 text documents, each with a comment about a piece of work a student submitted. Each text document is classified as positive “1” or negative “0”. This is a simple sentiment analysis problem.

First, we will define the documents and their class labels.

```
1 # define documents
2 docs = ['Well done!',
3         'Good work',
4         'Great effort',
5         'nice work',
6         'Excellent!',
7         'Weak',
8         'Poor effort!',
9         'not good',
10        'poor work',
11        'Could have done better.']
12 # define class labels
13 labels = array([1,1,1,1,1,0,0,0,0,0])
```

Next, we can integer encode each document. This means we convert each document into a sequence of integers. We could experiment with other encodings like counts or TF-IDF.

Keras provides the `one_hot()` function that creates a one-hot encoding.

We will estimate the vocabulary size of 50, which is much larger than needed to reduce the probability of collisions from the hash function.

```
1 # integer encode the documents
2 vocab_size = 50
3 encoded_docs = [one_hot(d, vocab_size) for d in docs]
4 print(encoded_docs)
```

The sequences have different lengths and Keras prefers inputs to be vectorized and all inputs to have the same length. We will pad all input sequences to have the length of 4. Again, we can do this with a built in Keras function, in this case the `pad_sequences()` function.

```
1 # pad documents to a max length of 4 words
2 max_length = 4
3 padded_docs = pad_sequences(encoded_docs, maxlen=max_length, padding='post')
4 print(padded_docs)
```

We are now ready to define our *Embedding* layer as part of our neural network model.

The *Embedding* has a vocabulary of 50 and an input length of 4. We will choose a small embedding space of 8 dimensions.

The model is a simple binary classification model. Importantly, the output from the *Embedding* layer will be 4 vectors of 8 dimensions each, one for each word. We flatten this to a one 32-element vector to pass on to the *Dense* output layer.

## Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**. Find out how in this *free* and *practical* course.

START MY EMAIL COURSE

Start Machine Learning

```

1 # define the model
2 model = Sequential()
3 model.add(Embedding(vocab_size, 8, input_length=max_length))
4 model.add(Flatten())
5 model.add(Dense(1, activation='sigmoid'))
6 # compile the model
7 model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
8 # summarize the model
9 print(model.summary())

```

Finally, we can fit and evaluate the classification model.

```

1 # fit the model
2 model.fit(padded_docs, labels, epochs=50, verbose=0)
3 # evaluate the model
4 loss, accuracy = model.evaluate(padded_docs, labels, verbose=0)
5 print('Accuracy: %f' % (accuracy*100))

```

The complete code listing is provided below.

```

1 from numpy import array
2 from keras.preprocessing.text import one_hot
3 from keras.preprocessing.sequence import pad_sequences
4 from keras.models import Sequential
5 from keras.layers import Dense
6 from keras.layers import Flatten
7 from keras.layers.embeddings import Embedding
8 # define documents
9 docs = ['Well done!',
10         'Good work',
11         'Great effort',
12         'nice work',
13         'Excellent!',
14         'Weak',
15         'Poor effort!',
16         'not good',
17         'poor work',
18         'Could have done better.']
19 # define class labels
20 labels = array([1,1,1,1,1,0,0,0,0,0])
21 # integer encode the documents
22 vocab_size = 50
23 encoded_docs = [one_hot(d, vocab_size) for d in docs]
24 print(encoded_docs)
25 # pad documents to a max length of 4 words
26 max_length = 4
27 padded_docs = pad_sequences(encoded_docs, maxlen=max_length, padding='post')
28 print(padded_docs)
29 # define the model
30 model = Sequential()
31 model.add(Embedding(vocab_size, 8, input_length=max_length))
32 model.add(Flatten())
33 model.add(Dense(1, activation='sigmoid'))
34 # compile the model
35 model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
36 # summarize the model
37 print(model.summary())
38 # fit the model
39 model.fit(padded_docs, labels, epochs=50, verbose=0)
40 # evaluate the model
41 loss, accuracy = model.evaluate(padded_docs, labels, verbose=0)
42 print('Accuracy: %f' % (accuracy*100))

```

## Start Machine Learning

You can master applied Machine Learning  
**without math or fancy degrees.**  
Find out how in this *free* and *practical* course.

START MY EMAIL COURSE

Start Machine Learning

Running the example first prints the integer encoded documents.

```
1 [[6, 16], [42, 24], [2, 17], [42, 24], [18], [17], [22, 17], [27, 42], [22, 24], [49, 46, 16, 34]]
```

Then the padded versions of each document are printed, making them all uniform length.

```
1 [[ 6 16  0  0]
2  [42 24  0  0]
3  [ 2 17  0  0]
4  [42 24  0  0]
5  [18  0  0  0]
6  [17  0  0  0]
7  [22 17  0  0]
8  [27 42  0  0]
9  [22 24  0  0]
10 [49 46 16 34]]
```

After the network is defined, a summary of the layers is printed. The Embedding layer is a 4x8 matrix and this is squared to get the final output shape of 32.

```
1 -----
2 Layer (type)                Output Shape
3 -----
4 embedding_1 (Embedding)      (None, 4, 8)
5 -----
6 flatten_1 (Flatten)          (None, 32)
7 -----
8 dense_1 (Dense)              (None, 1)
9 -----
10 Total params: 433
11 Trainable params: 433
12 Non-trainable params: 0
13 -----
```

Start Machine Learning

You can master applied Machine Learning without math or fancy degrees. Find out how in this free and practical course.

Email Address

START MY EMAIL COURSE

Finally, the accuracy of the trained model is printed, showing that it learned the training dataset perfectly (which is not surprising).

```
1 Accuracy: 100.000000
```

You could save the learned weights from the Embedding layer to file for later use in other models.

You could also use this model generally to classify other documents that have the same kind vocabulary seen in the test dataset.

Next, let's look at loading a pre-trained word embedding in Keras.

## 4. Example of Using Pre-Trained GloVe Embedding

The Keras Embedding layer can also use a word embedding learned elsewhere.

It is common in the field of Natural Language Processing to learn, save, and make freely available word embeddings.

For example, the researchers behind GloVe method provide a suite of pre-trained word embeddings on their website released under a public domain license.

- **GloVe: Global Vectors for Word Representation**

The smallest package of embeddings is 822Mb, called “*glove.6B.zip*”. It was trained on a dataset of one billion tokens (words) with a vocabulary of 400 thousand words. There are a few different embedding vector sizes, including 50, 100, 200 and 300 dimensions.

You can download this collection of embeddings and we can seed the Keras *Embedding* layer with weights from the pre-trained embedding for the words in your training dataset.

This example is inspired by an example in the Keras project: [pretrained\\_word\\_embeddings.py](#).

After downloading and unzipping, you will see a few files, one of which is “*glove.6B.100d.txt*”, which contains a 100-dimensional version of the embedding

If you peek inside the file, you will see a token (word) followed by its embedding vector. For example, below are the first line of the embedding

```
1 the -0.038194 -0.24487 0.72812 -0.39961 0.0831
```

As in the previous section, the first step is to define the sequences to be the same length.

In this case, we need to be able to map words to integers

Keras provides a `Tokenizer` class that can be fit on the texts consistently by calling the `texts_to_sequences()` method on the `Tokenizer` class, and provides access to the dictionary mapping of words to integers in a `word_index` attribute.

```
1 # define documents
2 docs = ['Well done!',
3         'Good work',
4         'Great effort',
5         'nice work',
6         'Excellent!',
7         'Weak',
8         'Poor effort!',
9         'not good',
10        'poor work',
11        'Could have done better.']
12 # define class labels
13 labels = array([1,1,1,1,1,0,0,0,0,0])
14 # prepare tokenizer
15 t = Tokenizer()
16 t.fit_on_texts(docs)
17 vocab_size = len(t.word_index) + 1
18 # integer encode the documents
19 encoded_docs = t.texts_to_sequences(docs)
20 print(encoded_docs)
21 # pad documents to a max length of 4 words
22 max_length = 4
23 padded_docs = pad_sequences(encoded_docs, maxlen=max_length, padding='post')
24 print(padded_docs)
```

Next, we need to load the entire GloVe word embedding array.

## Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**. Find out how in this *free* and *practical* course.

START MY EMAIL COURSE

Start Machine Learning

```

1 # load the whole embedding into memory
2 embeddings_index = dict()
3 f = open('glove.6B.100d.txt')
4 for line in f:
5     values = line.split()
6     word = values[0]
7     coefs = asarray(values[1:], dtype='float32')
8     embeddings_index[word] = coefs
9 f.close()
10 print('Loaded %s word vectors.' % len(embeddings_index))

```

This is pretty slow. It might be better to filter the embedding for the unique words in your training data.

Next, we need to create a matrix of one embedding for each word in the training dataset. We can do that by enumerating all unique words in the `Tokenizer.word_index` and locating the embedding weight vector from the loaded GloVe embedding.

The result is a matrix of weights only for words we will

```

1 # create a weight matrix for words in training
2 embedding_matrix = zeros((vocab_size, 100))
3 for word, i in t.word_index.items():
4     embedding_vector = embeddings_index.get(word)
5     if embedding_vector is not None:
6         embedding_matrix[i] = embedding_vector

```

Now we can define our model, fit, and evaluate it as b

The key difference is that the embedding layer can be chose the 100-dimensional version, therefore the Embedding layer must be defined with `output_dim` set to 100. Finally, we do not want to update the learned word weights in this model, therefore we will set the `trainable` attribute for the model to be `False`.

```

1 e = Embedding(vocab_size, 100, weights=[embedding_matrix], input_length=4, trainable=False)

```

The complete worked example is listed below.

```

1 from numpy import array
2 from numpy import asarray
3 from numpy import zeros
4 from keras.preprocessing.text import Tokenizer
5 from keras.preprocessing.sequence import pad_sequences
6 from keras.models import Sequential
7 from keras.layers import Dense
8 from keras.layers import Flatten
9 from keras.layers import Embedding
10 # define documents
11 docs = ['Well done!',
12         'Good work',
13         'Great effort',
14         'nice work',
15         'Excellent!',
16         'Weak',
17         'Poor effort!',
18         'not good',
19         'poor work',
20         'Could have done better.']
21 # define class labels
22 labels = array([1,1,1,1,1,0,0,0,0,0])

```

## Start Machine Learning

You can master applied Machine Learning  
without math or fancy degrees.  
Find out how in this free and practical course.

START MY EMAIL COURSE

We

Start Machine Learning



```

23 # prepare tokenizer
24 t = Tokenizer()
25 t.fit_on_texts(docs)
26 vocab_size = len(t.word_index) + 1
27 # integer encode the documents
28 encoded_docs = t.texts_to_sequences(docs)
29 print(encoded_docs)
30 # pad documents to a max length of 4 words
31 max_length = 4
32 padded_docs = pad_sequences(encoded_docs, maxlen=max_length, padding='post')
33 print(padded_docs)
34 # load the whole embedding into memory
35 embeddings_index = dict()
36 f = open('../glove_data/glove.6B/glove.6B.100d.txt')
37 for line in f:
38     values = line.split()
39     word = values[0]
40     coefs = asarray(values[1:], dtype='float32')
41     embeddings_index[word] = coefs
42 f.close()
43 print('Loaded %s word vectors.' % len(embeddings_index))
44 # create a weight matrix for words in training
45 embedding_matrix = zeros((vocab_size, 100))
46 for word, i in t.word_index.items():
47     embedding_vector = embeddings_index.get(word)
48     if embedding_vector is not None:
49         embedding_matrix[i] = embedding_vector
50 # define model
51 model = Sequential()
52 e = Embedding(vocab_size, 100, weights=[embedding_matrix])
53 model.add(e)
54 model.add(Flatten())
55 model.add(Dense(1, activation='sigmoid'))
56 # compile the model
57 model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
58 # summarize the model
59 print(model.summary())
60 # fit the model
61 model.fit(padded_docs, labels, epochs=50, verbose=0)
62 # evaluate the model
63 loss, accuracy = model.evaluate(padded_docs, labels, verbose=0)
64 print('Accuracy: %f' % (accuracy*100))

```

## Start Machine Learning

You can master applied Machine Learning  
without math or fancy degrees.  
Find out how in this *free* and *practical* course.

START MY EMAIL COURSE

Running the example may take a bit longer, but then demonstrates that it is just as capable of fitting this simple problem.

```

1 [[6, 2], [3, 1], [7, 4], [8, 1], [9], [10], [5, 4], [11, 3], [5, 1], [12, 13, 2, 14]]
2
3 [[ 6  2  0  0]
4  [ 3  1  0  0]
5  [ 7  4  0  0]
6  [ 8  1  0  0]
7  [ 9  0  0  0]
8  [10  0  0  0]
9  [ 5  4  0  0]
10 [11  3  0  0]
11 [ 5  1  0  0]
12 [12 13  2 14]]
13
14 Loaded 400000 word vectors.
15
16
17 Layer (type)                Output Shape
18 =====

```

Start Machine Learning

```

19 embedding_1 (Embedding)      (None, 4, 100)      1500
20 -----
21 flatten_1 (Flatten)         (None, 400)          0
22 -----
23 dense_1 (Dense)              (None, 1)            401
24 =====
25 Total params: 1,901
26 Trainable params: 401
27 Non-trainable params: 1,500
28 -----
29
30
31 Accuracy: 100.000000

```

In practice, I would encourage you to experiment with learning a word embedding using a pre-trained embedding that is fixed and trying to perform learning on top of a pre-trained embedding.

See what works best for your specific problem.

## Further Reading

This section provides more resources on the topic if you are interested.

- [Word Embedding on Wikipedia](#)
- [Keras Embedding Layer API](#)
- [Using pre-trained word embeddings in a Keras model](#)
- [Example of using a pre-trained GloVe Embedding](#)
- [GloVe Embedding](#)
- [An overview of word embeddings and their connection to distributional semantic models, 2016](#)
- [Deep Learning, NLP, and Representations, 2014](#)

## Summary

In this tutorial, you discovered how to use word embeddings for deep learning in Python with Keras.

Specifically, you learned:

- About word embeddings and that Keras supports word embeddings via the Embedding layer.
- How to learn a word embedding while fitting a neural network.
- How to use a pre-trained word embedding in a neural network.

Do you have any questions?

Ask your questions in the comments below and I will do my best to answer.

### Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**. Find out how in this *free* and *practical* course.


START MY EMAIL COURSE

## Develop Deep Learning models for Text Data Today!

Develop Your Own Text models in Minutes

...with just a few lines of code

Start Machine Learning



Discover how in my new Ebook:  
**Deep Learning for Natural Language Processing**

It provides **self-study tutorials** on topics like:  
*Bag-of-Words, Word Embedding, Language Models, Caption Generation, Text Translation* and much more...

## Finally Bring Deep Learning to your Natural Language Processing Projects

Skip the Academics. Just Results.

SEE WHAT'S INSIDE

### Start Machine Learning

You can master applied Machine Learning  
**without math or fancy degrees.**  
Find out how in this *free* and *practical* course.

START MY EMAIL COURSE

Tweet

Share

Share



### About Jason Brownlee

Jason Brownlee, PhD is a machine learning expert who teaches modern machine learning methods via hands-on tutorials.  
[View all posts by Jason Brownlee](#) →

< [How to Prepare Text Data for Deep Learning with Keras](#)

[How to Develop Word Embeddings in Python with Gensim](#) >

## 429 Responses to *How to Use Word Embedding Layers for Deep Learning with Keras*



**Mohammad** October 4, 2017 at 7:58 am #

REPLY ↩

Thank you Jason,  
I am excited to read more NLP posts.



**Jason Brownlee** October 4, 2017 at 8:03 am #

REPLY ↩

Thanks.

Start Machine Learning



**sherry** July 22, 2019 at 7:22 pm #

REPLY ↩

after embedding, have to have a "Flatten()" layer? In my project, I used a dense layer directly after embedding. is it ok?



**Jason Brownlee** July 23, 2019 at 7:59 am #

REPLY ↩

Try it and see.



**Peter Nduru** October 8, 2019 at 6:06 am #

I appreciate how well updated you keep the start reading is the update date. thank you very much

## Start Machine Learning



You can master applied Machine Learning **without math or fancy degrees**. Find out how in this *free* and *practical* course.

START MY EMAIL COURSE



**Jason Brownlee** October 8, 2019 at 8:10 am #

You're welcome.

I require all of the code to work and keep working!



**Martin** October 11, 2019 at 4:09 am #

REPLY ↩

Hi, Jason:

when one\_hot encoding is used, why is padding necessary? Doesn't one\_hot encoding already create an input of equal length?



**Jason Brownlee** October 11, 2019 at 6:25 am #

REPLY ↩

The one hot encoding is for one variable at one time step, e.g. features.

Padding is needed to make all sequences have the same number of time steps.

See this:

<https://machinelearningmastery.com/faq/single-faq/what-is-the-difference-between-samples-timesteps-and-features-for-lstm-input>

Start Machine Learning



**shiv** October 5, 2017 at 10:07 am #

REPLY ↩

I split my data into 80-20 test-train and I'm still getting 100% accuracy. Any idea why? It is ~99% on epoch 1 and the rest its 100%.



**Jason Brownlee** October 5, 2017 at 5:22 pm #

REPLY ↩

Consider using the procedure in this post to evaluate your model:  
<https://machinelearningmastery.com/evaluate-skill-deep-learning-models/>



**trulia** October 6, 2017 at 12:47 pm #

Use drop-out 20%, your model is overfit!!



**Sandy** October 6, 2017 at 2:44 pm #

Thank you Jason. I always find things easier  
I have a question about the vector of each word after training. For example, if I have a sentence "Well done!" will be represented in different vector from "Well done! Could have done better!" Is that right? I mean the presentation of each word will depend on the context of each sentence?

## Start Machine Learning

You can master applied Machine Learning  
**without math or fancy degrees.**  
Find out how in this *free* and *practical* course.

START MY EMAIL COURSE



**Jason Brownlee** October 7, 2017 at 5:48 am #

REPLY ↩

No, each word in the dictionary is represented differently, but the same word in different contexts will have the same representation.

It is the word in its different contexts that is used to define the representation of the word.

Does that help?



**Sandy** October 7, 2017 at 5:37 pm #

REPLY ↩

Yes, thank you. But I still have a question. We will train each context separately, then after training the first context, in this case is "Well done!", we will have a vector representation of the word "done". After training the second context, "Could have done better", we have another vector representation of the word "done". So, which vector will we choose to be the representation of the word "done"?

I might misunderstand the procedure of training

Start Machine Learning



**Jason Brownlee** October 8, 2017 at 8:32 am #

REPLY ↩

No. All examples where a word is used are used as part of the training of the representation of the word. There is only one representation for each word during and after training.



**Sandy** October 8, 2017 at 2:46 pm #

I got it. Thank you, Jason.



**Chiedu** October 7, 2017 at 5:36 pm #

Hi Jason,  
any ideas on how to “filter the embedding for the unique tutorial?”

## Start Machine Learning



You can master applied Machine Learning **without math or fancy degrees**. Find out how in this *free* and *practical* course.

START MY EMAIL COURSE



**Jason Brownlee** October 8, 2017 at 8:32 am

The mapping of word to vector dictionary is built into Gensim, you can access it directly to retrieve the representations for the words you want: `model.wv.vocab`



**mahna** April 28, 2018 at 2:31 am #

REPLY ↩

Hi Jason,  
I am really appreciated the time U spend to write this tutorial and also replying.  
My question is about “`model.wv.vocab`” you wrote. is it an address site?  
It does not work actually.



**Jason Brownlee** April 28, 2018 at 5:33 am #

REPLY ↩

No, it is an attribute on the model.



**Abbey** October 8, 2017 at 2:19 am #

REPLY ↩

Start Machine Learning

Hi, Jason

Good day.

I just need your suggestion and example. I have two different dataset, where one is structured and the other is unstructured. The goal is to use the structured to construct a representation for the unstructured, so apply use word embedding on the two input data but how can I find the average of the two embedding and flatten it to one before feeding the layer into CNN and LSTM.

Looking forward to your response.

Regards

Abbey



**Jason Brownlee** October 8, 2017 at 8:40 am

Sorry, what was your question?

If your question was if this is a good approach, my



**Abiodun Modupe** October 9, 2017 at 7:00 am

Hi, Jason

How can I find the average of the word embedding from the two input?

Regards

Abbey

## Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**. Find out how in this *free* and *practical* course.

START MY EMAIL COURSE



**Jason Brownlee** October 10, 2017 at 7:43 am #

REPLY ↩

Perhaps you could retrieve the vectors for each word and take their average?

Perhaps you can use the Gensim API to achieve this result?



**Rafael Sá** June 17, 2019 at 2:47 am #

REPLY ↩

Hi Jason,

I have a set of documents(1200 text of movie Scripts) and i want to use pretrained embeddings. But i want to update the vocabulary and train again adding the words of my corpus. Is that possible ?



**Jason Brownlee** June 17, 2019 at 8:00 am

Start Machine Learning